

# Performance Tuning For Azure Cosmos DB

Cloud Databases



Hasan Savran

SavranWeb Consulting



# Hasan Savran (He/him)

## Owner SavranWeb Consulting



- MS Data Platform MVP
- Azure Cosmos DB SME
- From Cleveland, USA
- 15+ years Web Development
- 8+ years Business Intelligence



<https://h-savran.blogspot.com/>



hasansavran



@savranweb

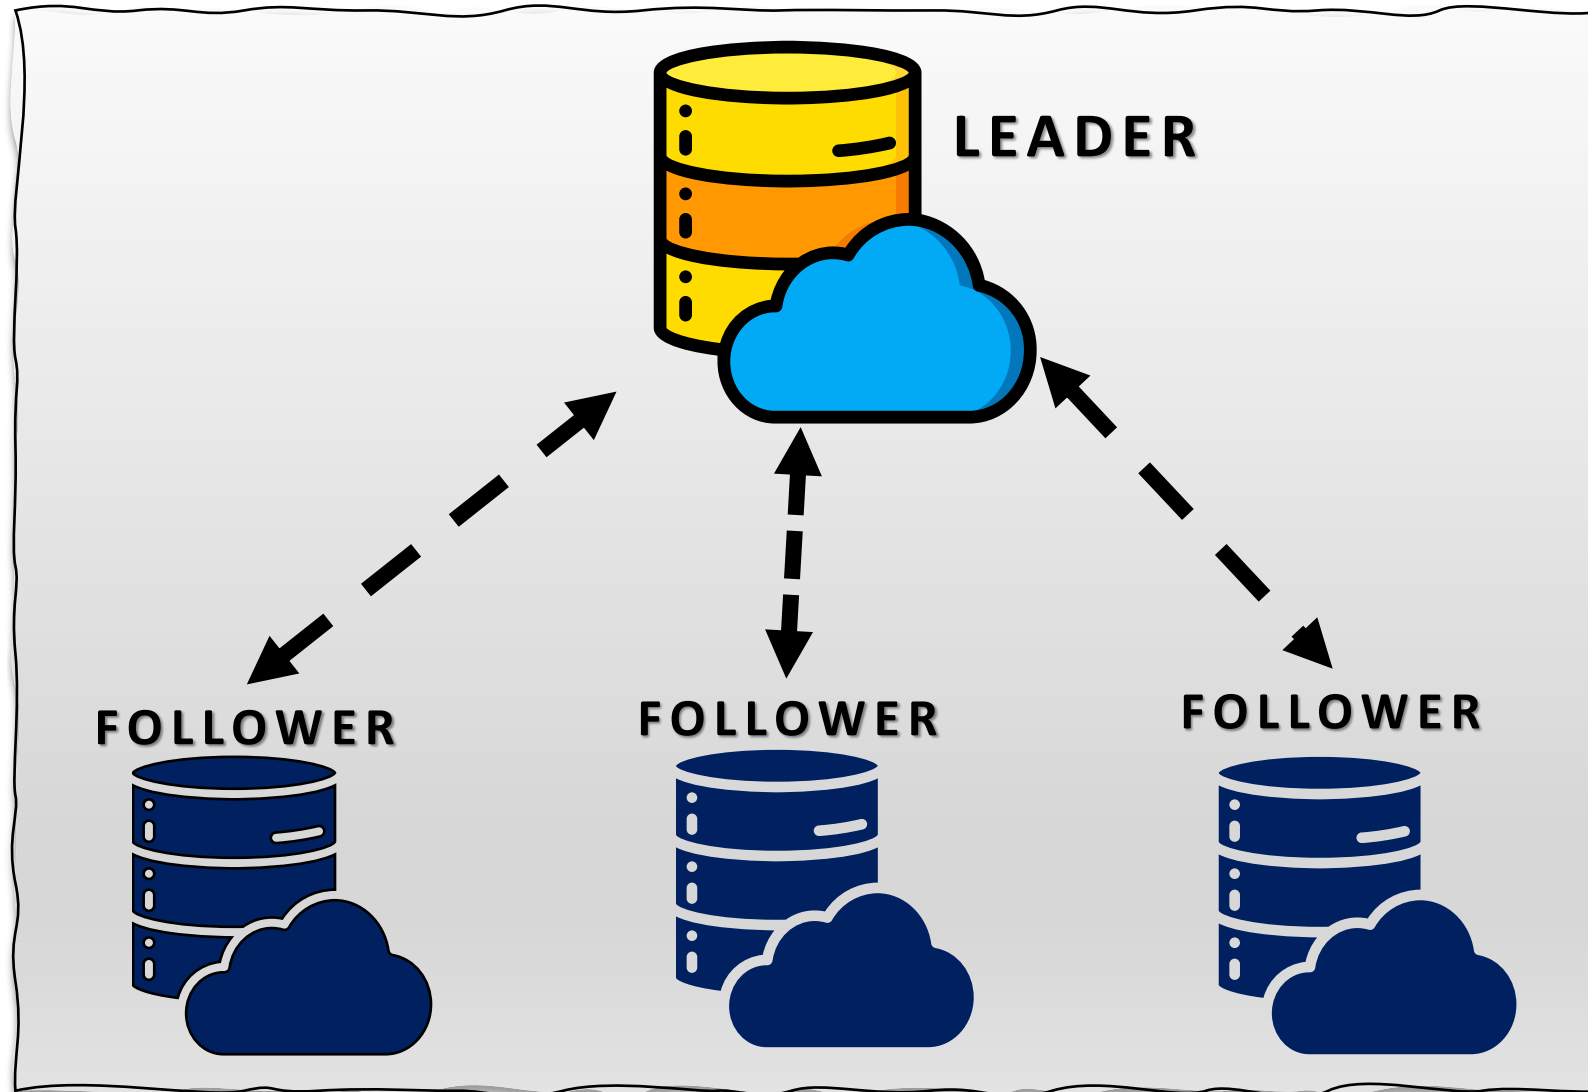


# COSMOS DB PHYSICAL PARTITION

## Limits

50 GB Data

10,000 RU



Query  
Engine

Storage  
Engine

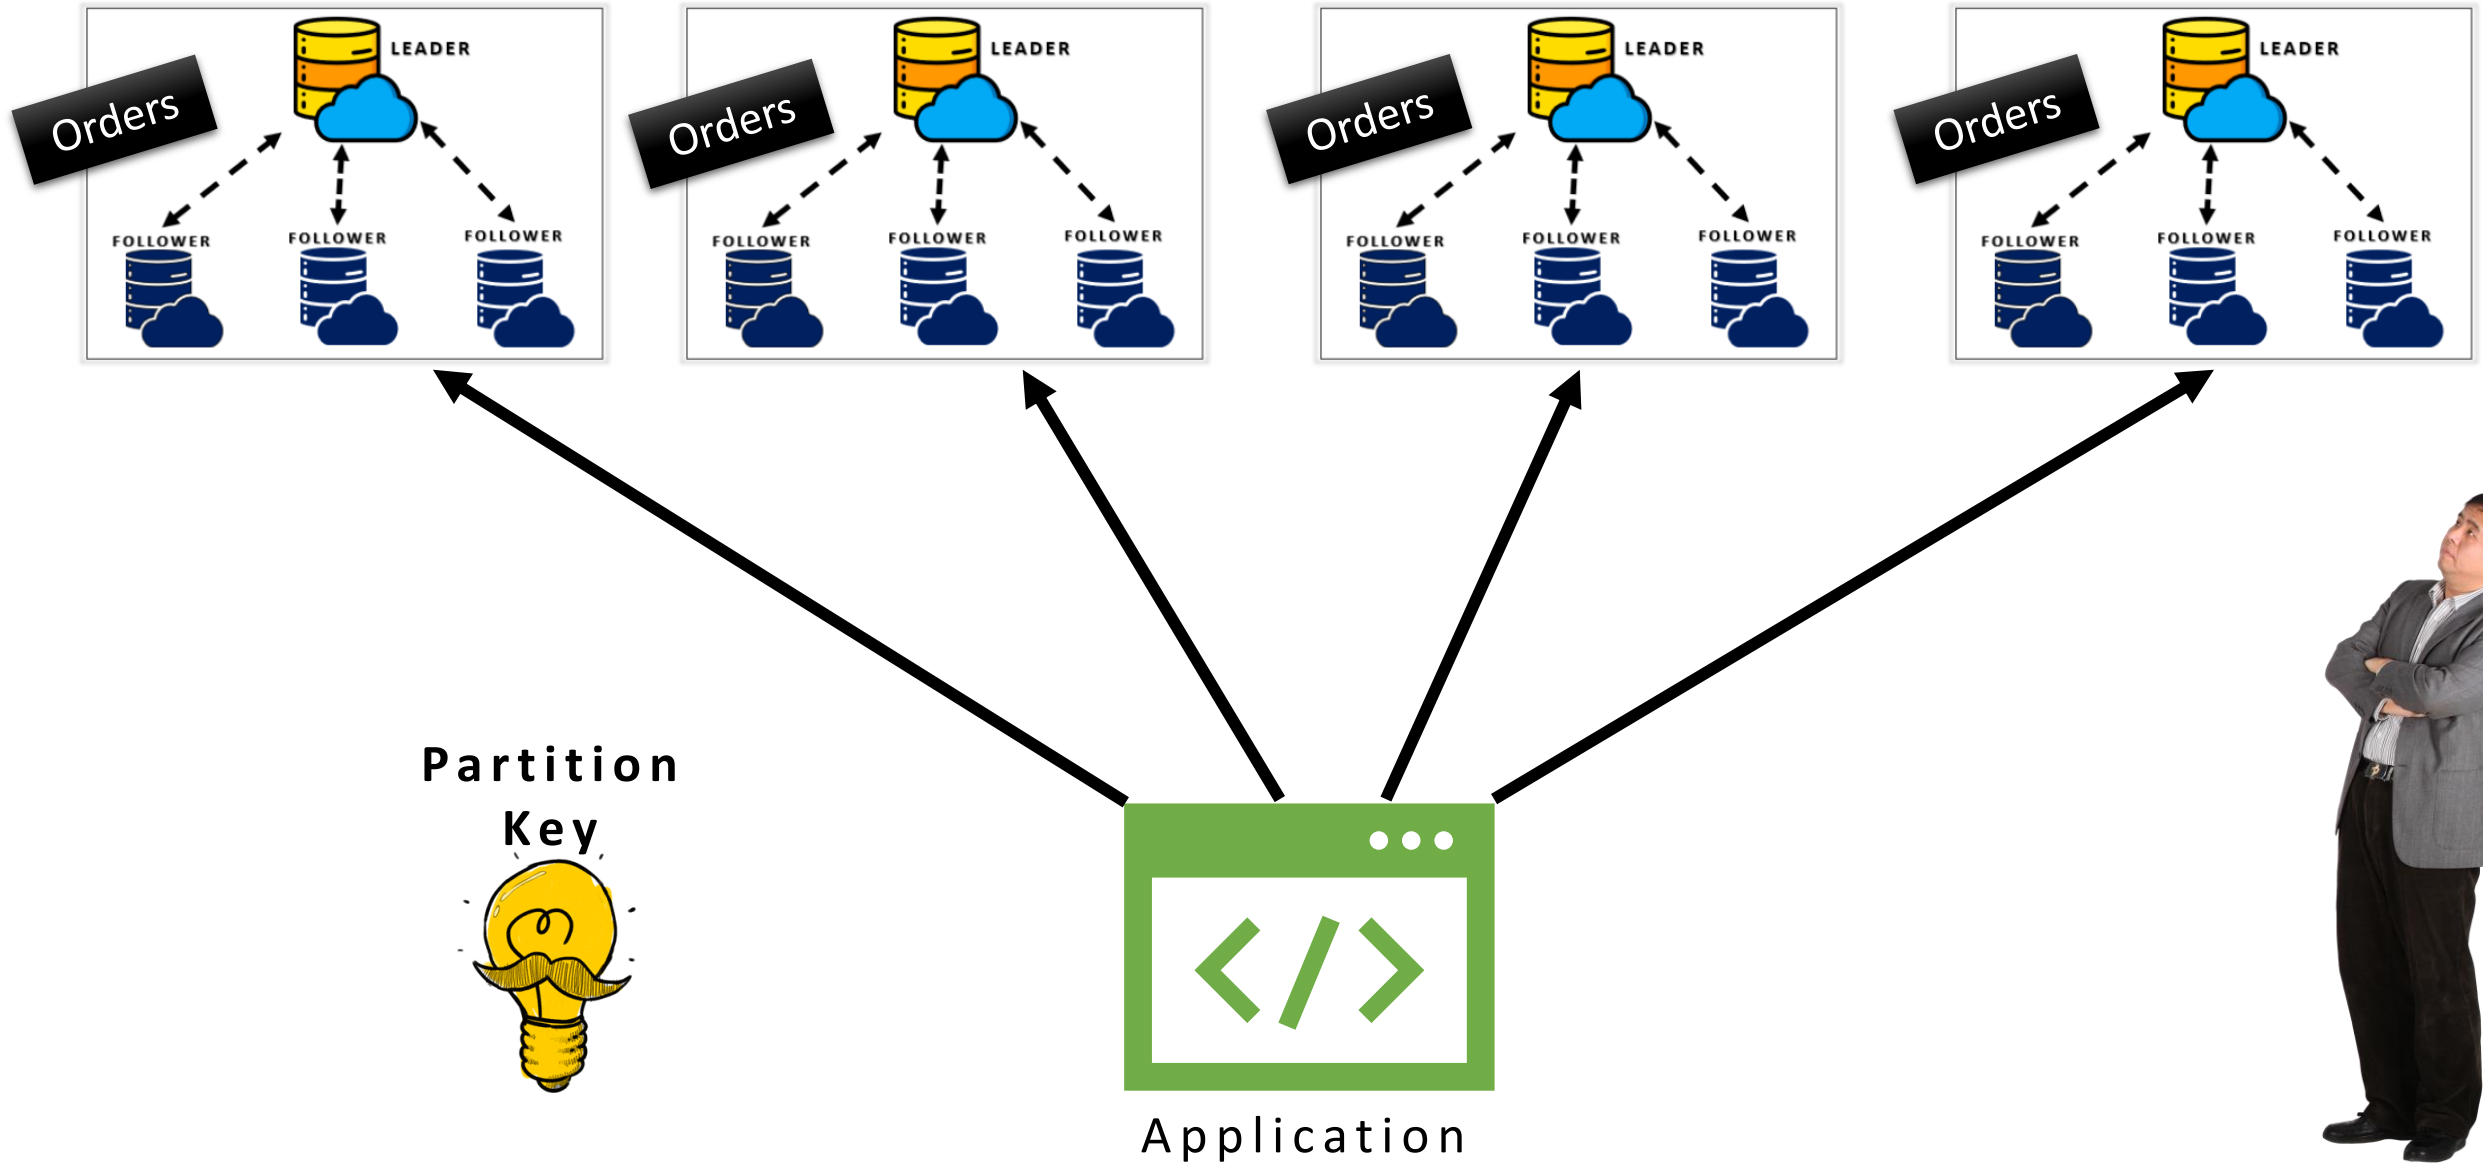
Indexing  
Engine

WITHOUT GLOBAL DISTRIBUTION

Database with 180 GB Data

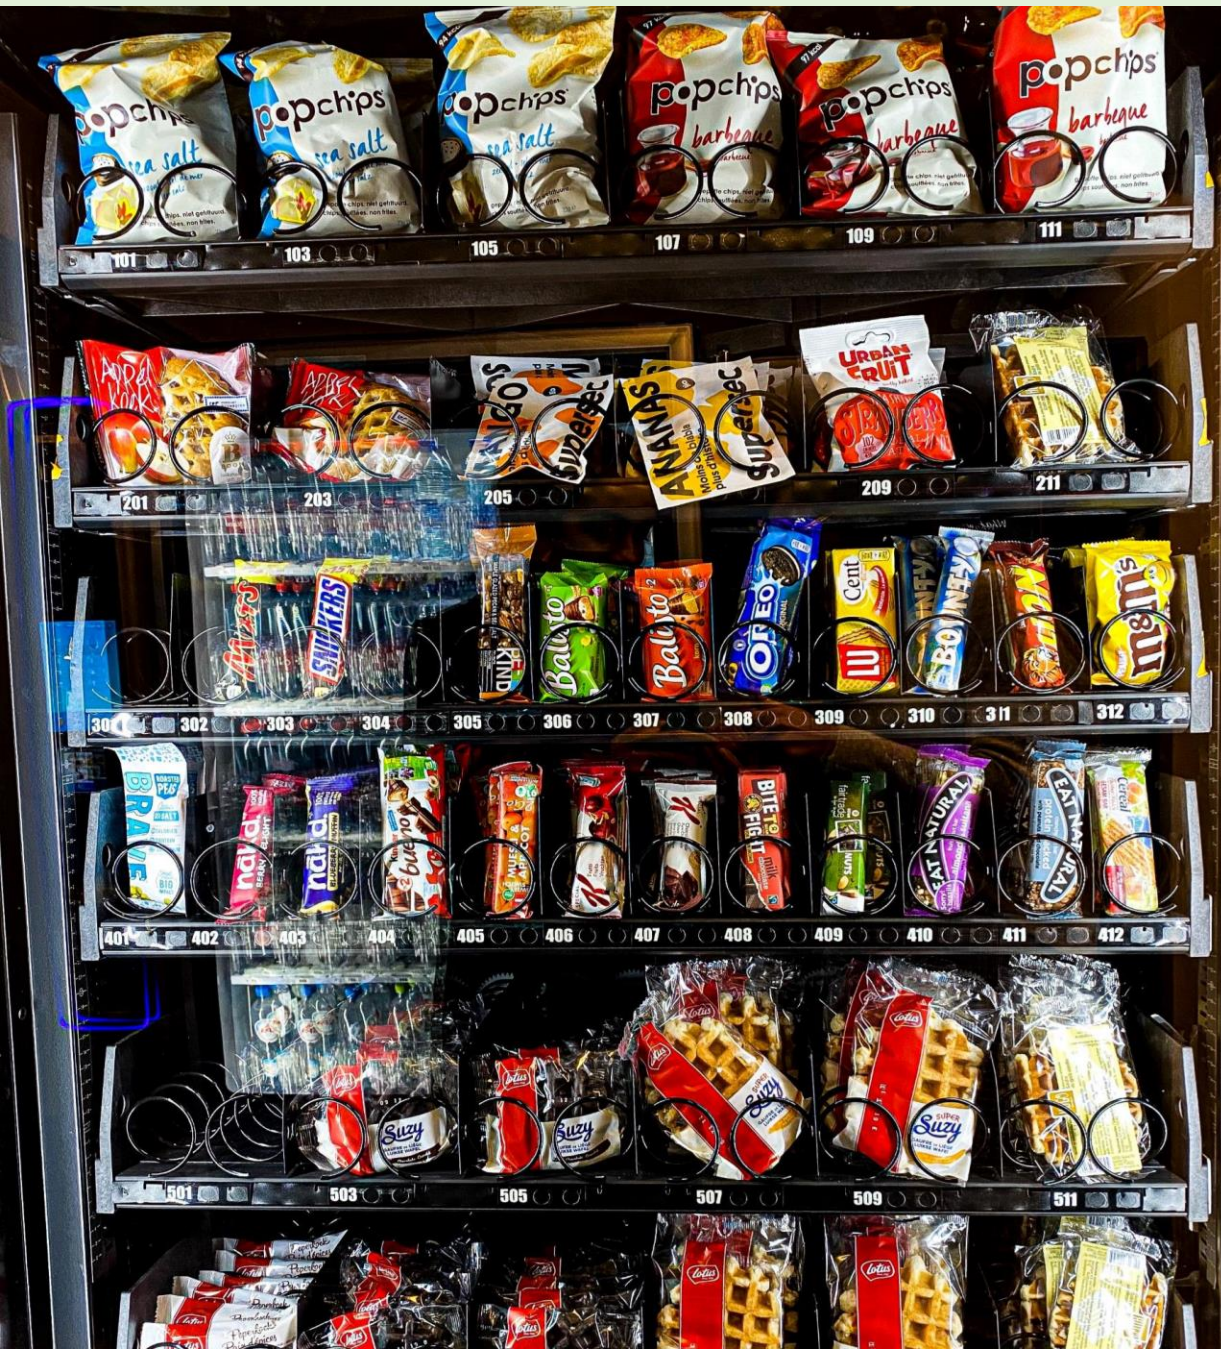
OR

Database with 20,000 Request Unit





## LOW RU & High Scalability



## HIGH RU & BIG MESS



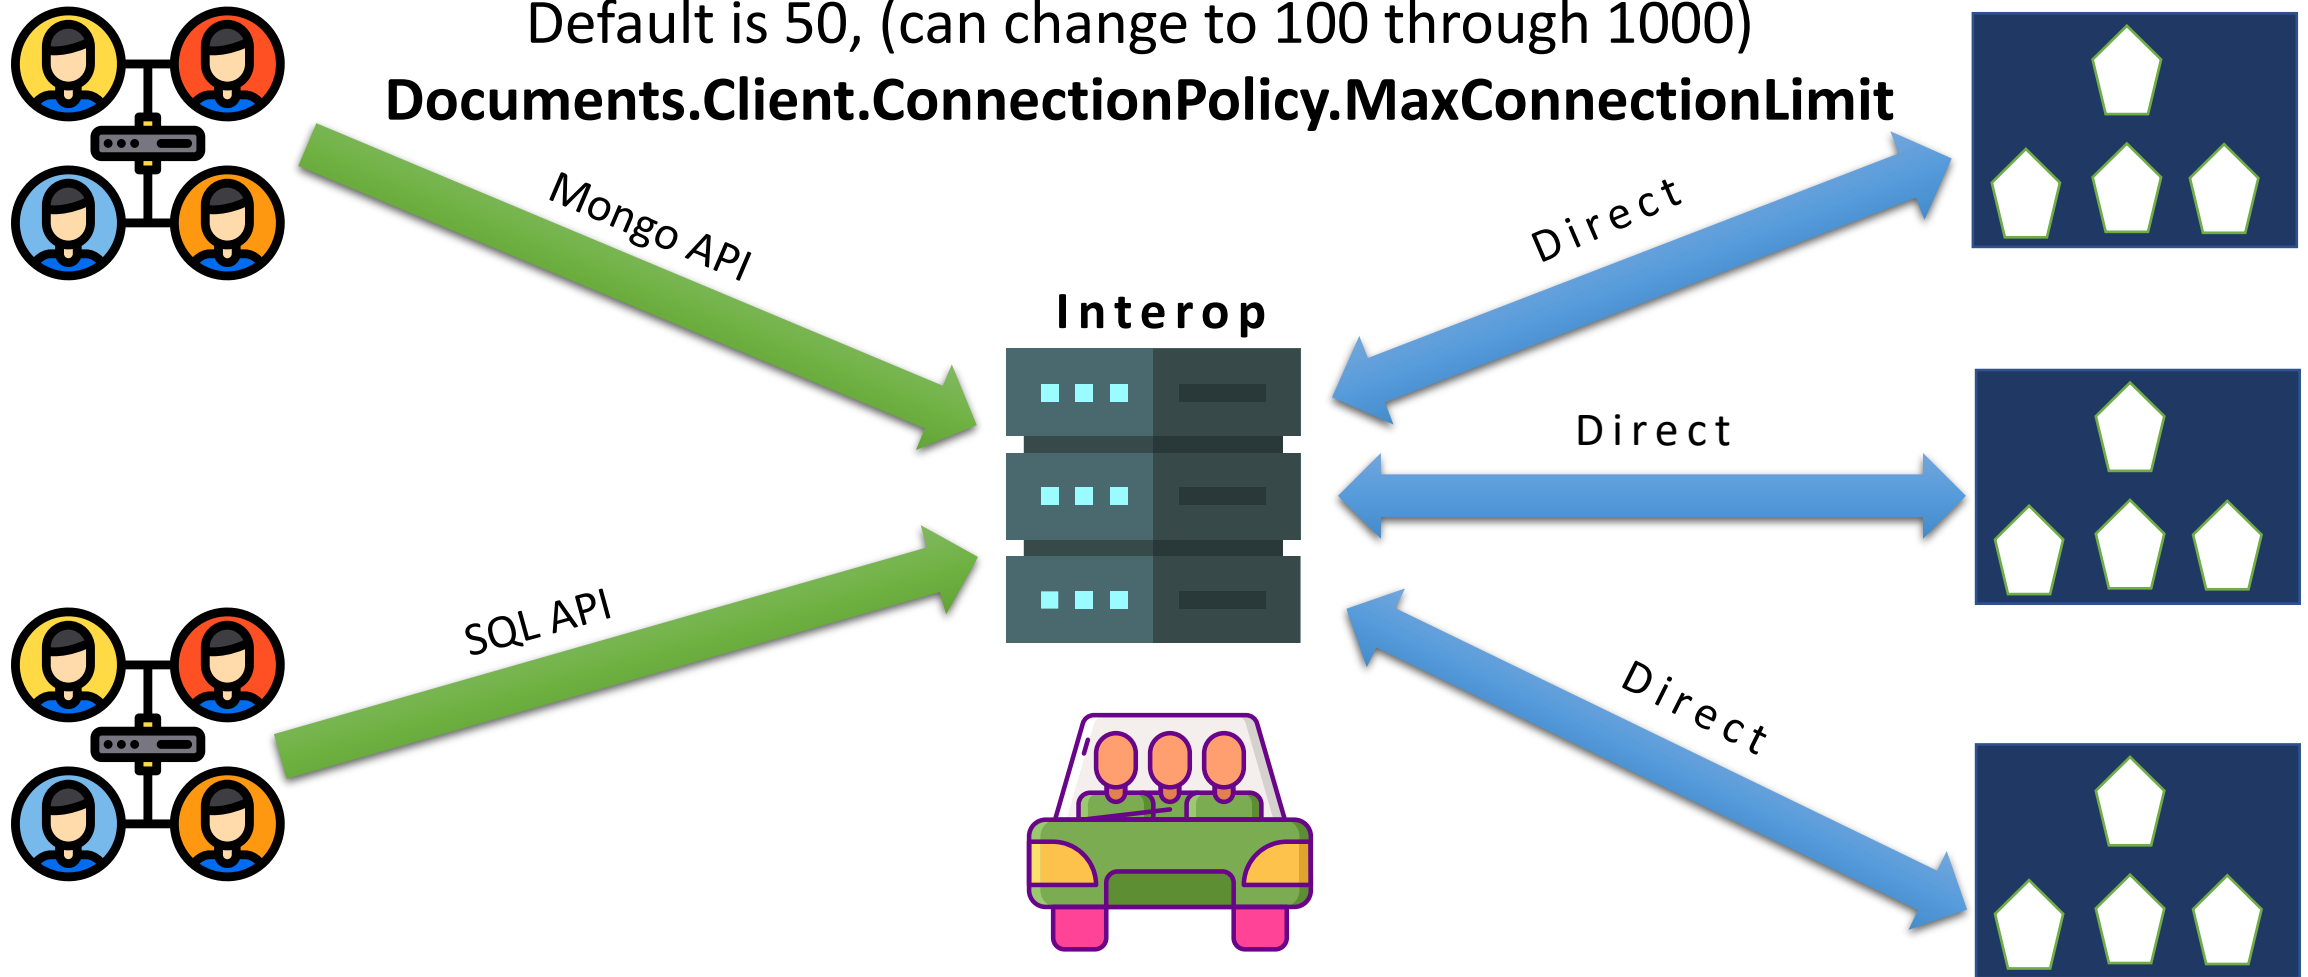


# Connecting to Cosmos DB

Increase **MaxConnections** per host/IP

Default is 50, (can change to 100 through 1000)

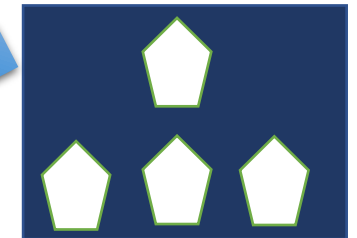
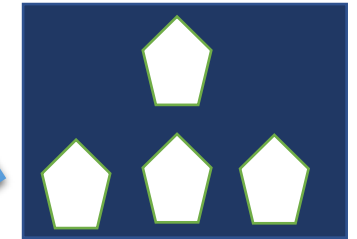
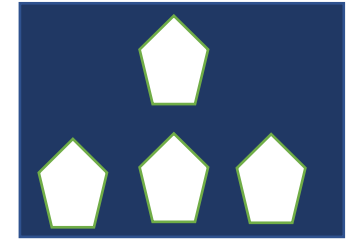
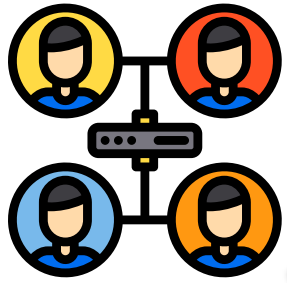
**Documents.Client.ConnectionPolicy.MaxConnectionLimit**



Gateway Mode uses **HTTPS & single DNS endpoint**

# Connecting to Cosmos DB

## Direct Connection(TCP)

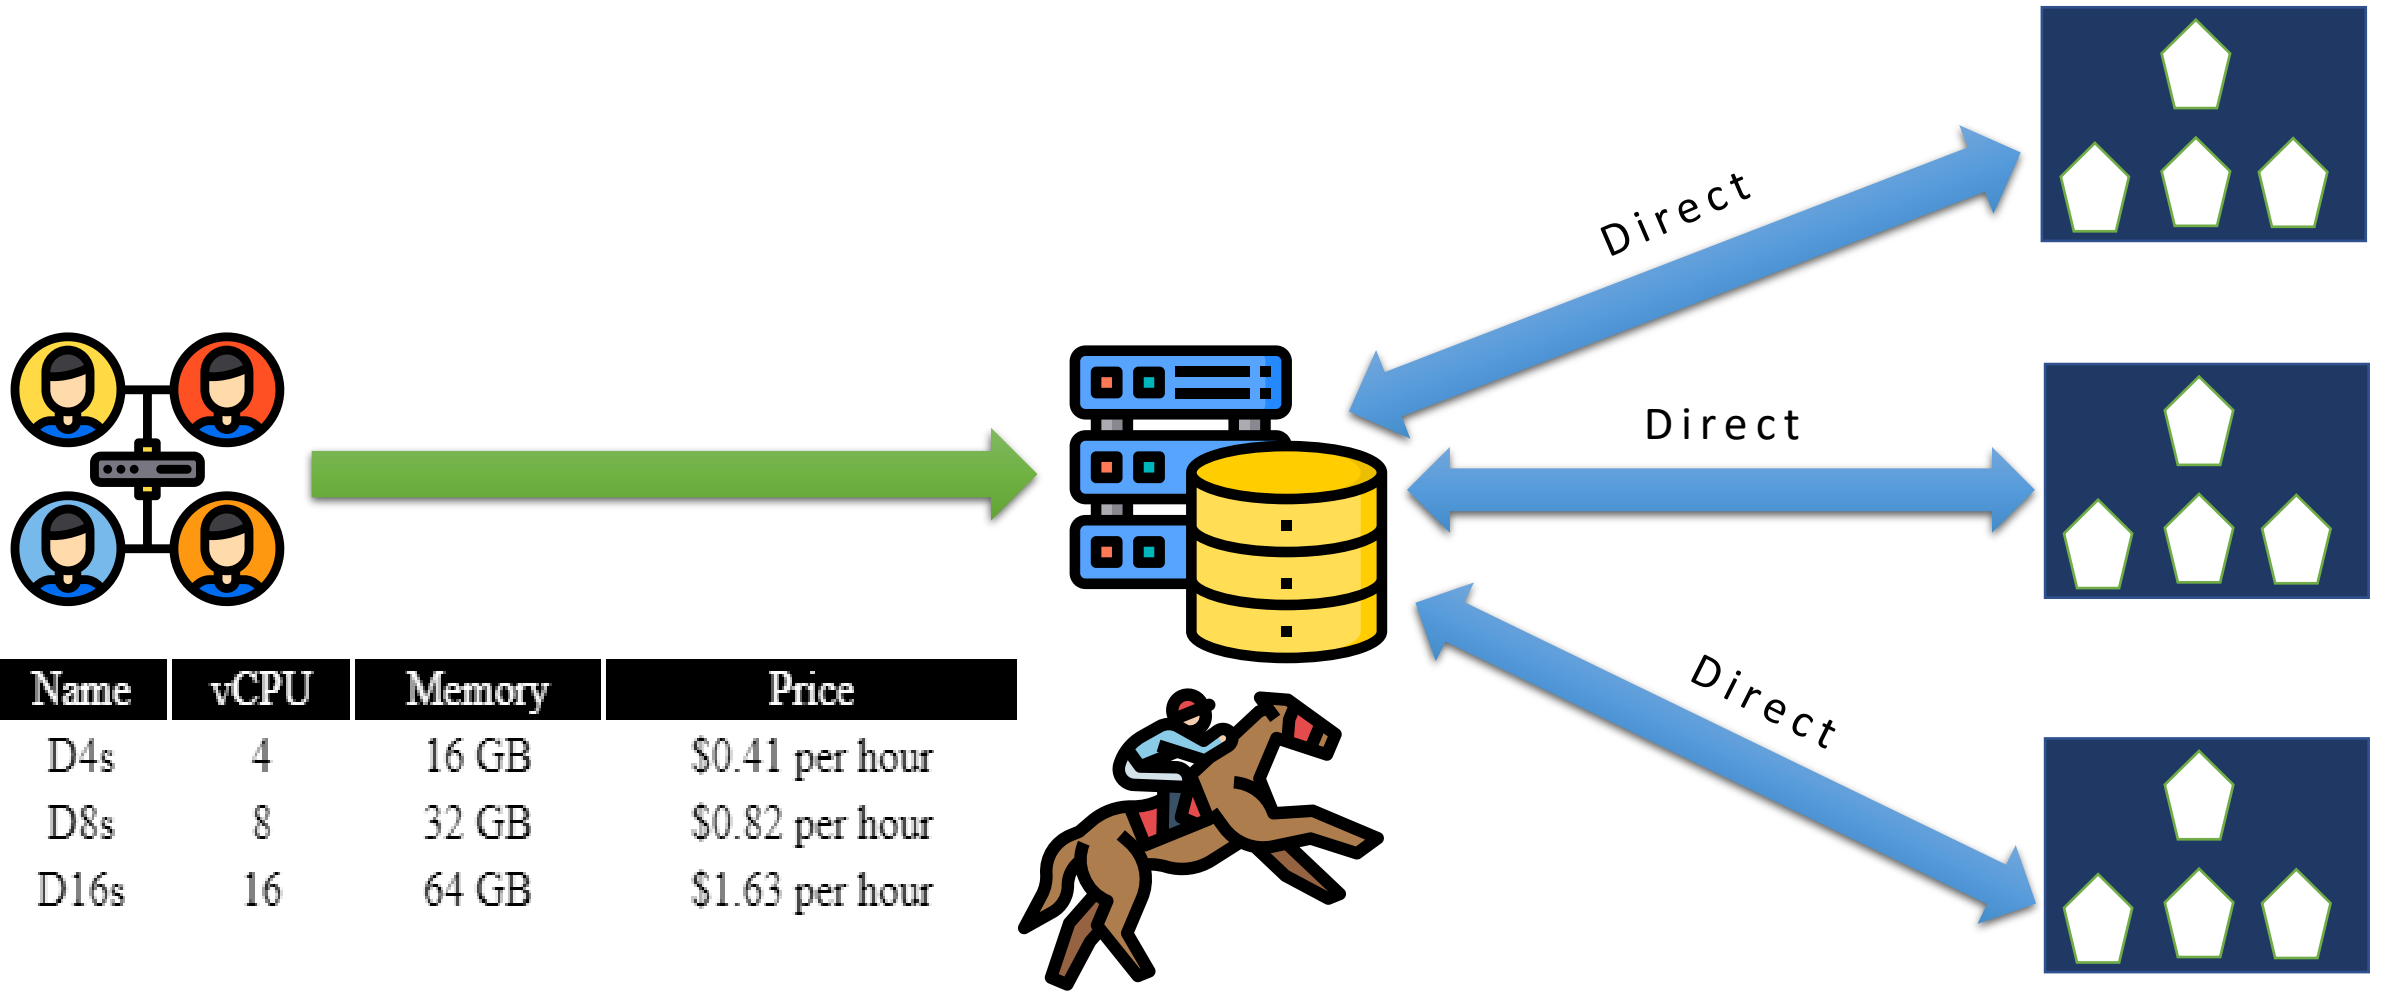


**.NET  
SDK**

**Java  
SDK**

```
var client = new CosmosClient(cstring, new CosmosClientOptions
{
    ...
    ConnectionMode = ConnectionMode.Direct
});
```

# Connecting to Cosmos DB

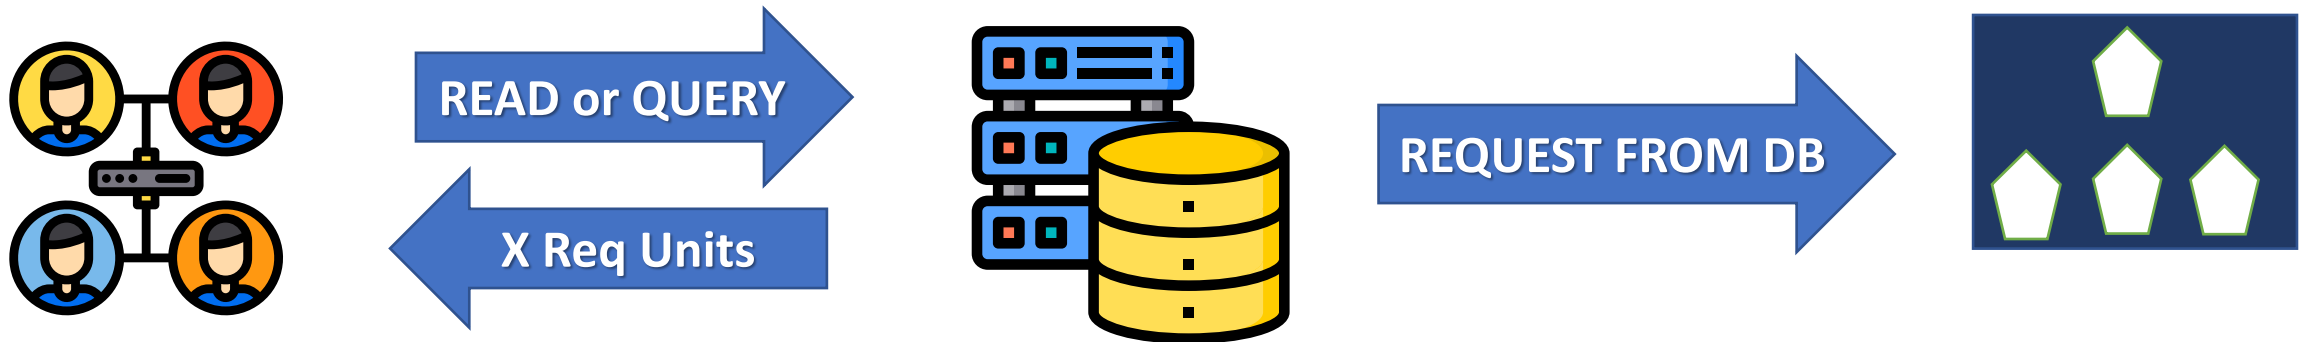


Dedicated Gateway Mode

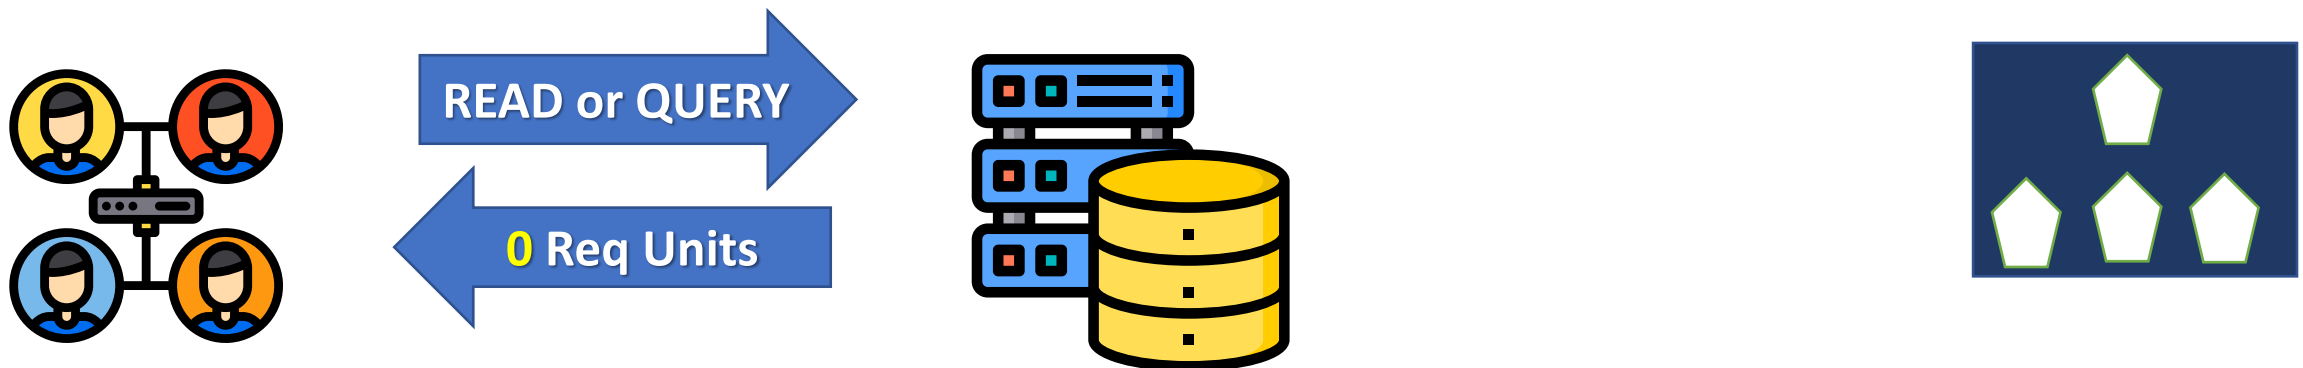


# Caching Data with Integrated Gateway

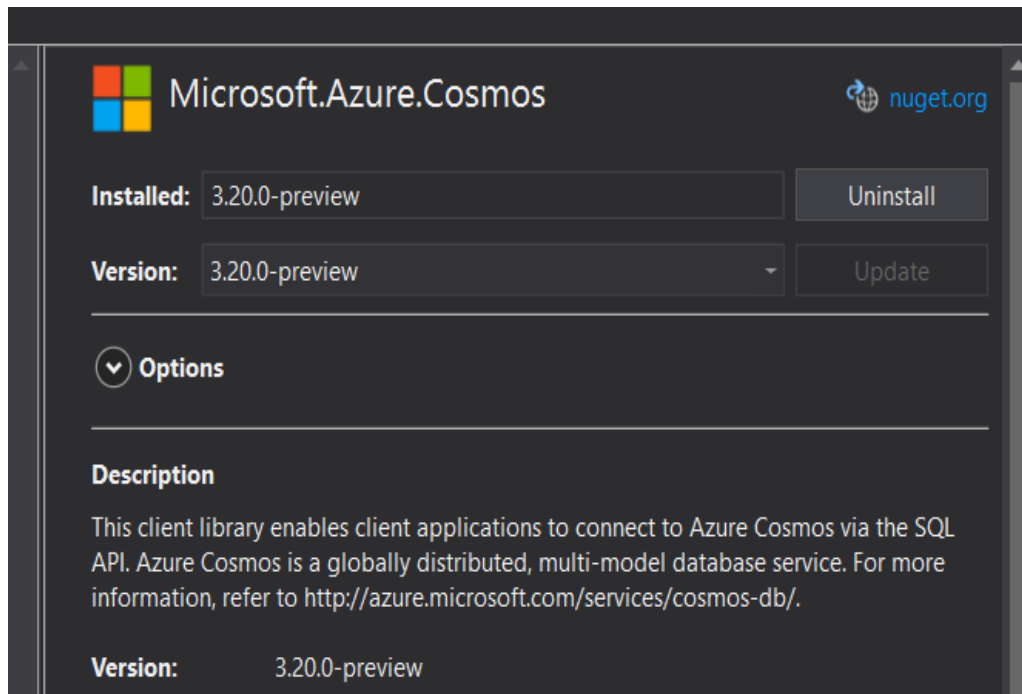
Consistency Level must be **EVENTUAL**



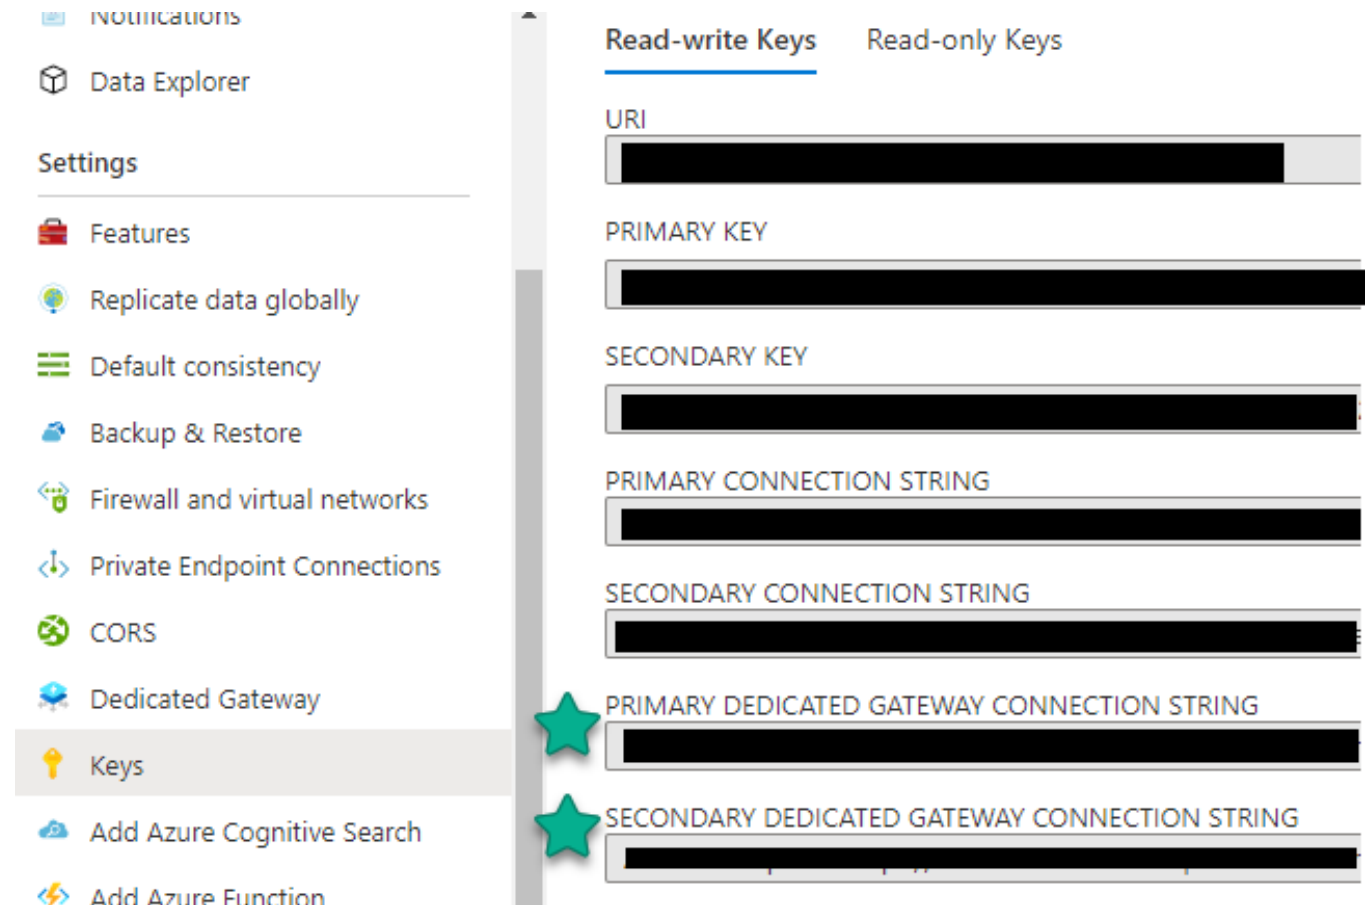
**SAME Request comes later**



# Caching Data with Integrated Gateway



**SDK 3.2.1 or later**



**Use Dedicated Gateway Connection String**

# Caching Data with Integrated Gateway

```
static async Task<List<StackOverflowPost>> TestCaching(int postId=0)
{
    var cosmosClient = new CosmosClient(connectionString,
        new CosmosClientOptions { ConnectionMode= ConnectionMode.Gateway});
    Container container = cosmosClient.GetContainer("Stackoverflow", "Posts");
    var cmd = "SELECT * FROM Posts o WHERE o.PostId < 500";
    var query = new QueryDefinition(cmd);
    var queryResultSetIterator = container.GetItemQueryIterator<StackOverflowPost>(query,
        requestOptions: new QueryRequestOptions
        {
            ConsistencyLevel = ConsistencyLevel.Eventual,
            DedicatedGatewayRequestOptions = new DedicatedGatewayRequestOptions
            {
                MaxIntegratedCacheStaleness = TimeSpan.FromMinutes(30)
            }
        });
    var posts = new List<StackOverflowPost>();
    double rq = 0;
    try
    {
        while (queryResultSetIterator.HasMoreResults)
        {
```

# Caching Data with Integrated Gateway

FIRST TIME

C:\> Select Microsoft Visual Studio Debug Console

```
Number of Posts : 592
```

```
Request Units : 32.25
```

```
C:\Program Files\dotnet\dotnet.exe (process 3004) exited  
To automatically close the console when debugging stops,  
please use the /console:None switch.  
Press any key to close this window . . .
```

LATER

C:\> Microsoft Visual Studio Debug Console

```
Number of Posts : 592
```

```
Request Units : 0
```

```
C:\Program Files\dotnet\dotnet.exe (process 22088) exited  
To automatically close the console when debugging stops,  
please use the /console:None switch.  
Press any key to close this window . . .
```



# SDK

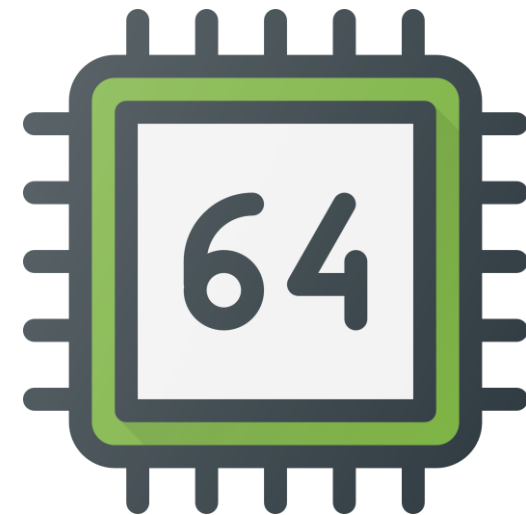
```
▼ queryInfo: {distinctType: 'None', top: 10, offset: nul...
  > aggregates: (0) []
    dCountInfo: null
    distinctType: 'None'
  > groupByAliases: (0) []
  > groupByAliasToAggregateType: {}
  > groupByExpressions: (0) []
    hasSelectValue: false
    limit: null
    offset: null
  > orderBy: (1) ['Ascending']
  > orderByExpressions: (1) ['c.ViewCount']
    rewrittenQuery: 'SELECT TOP 10 c._rid, [{"item": c.ViewCount}] AS orderByItems, c AS payload'
    top: 10
  > __proto__: Object
  > queryRanges: (1) [{...}]
  > __proto__: Object
  query: 'select top 10 * from c\norder by c.ViewCount'
```

```
SELECT TOP 10 c._rid,
[{"item": c.ViewCount}] AS orderByItems, c AS payload
FROM c
WHERE ({documentdb-formattableorderbyquery-filter})
ORDER BY c.ViewCount
```



# .NET & Hosting Recommendations

- Use 64-bit Windows host processing
- ServiceInterop.dll to parse and optimize queries locally for SQL SDK



*runtimeconfig.json* file:



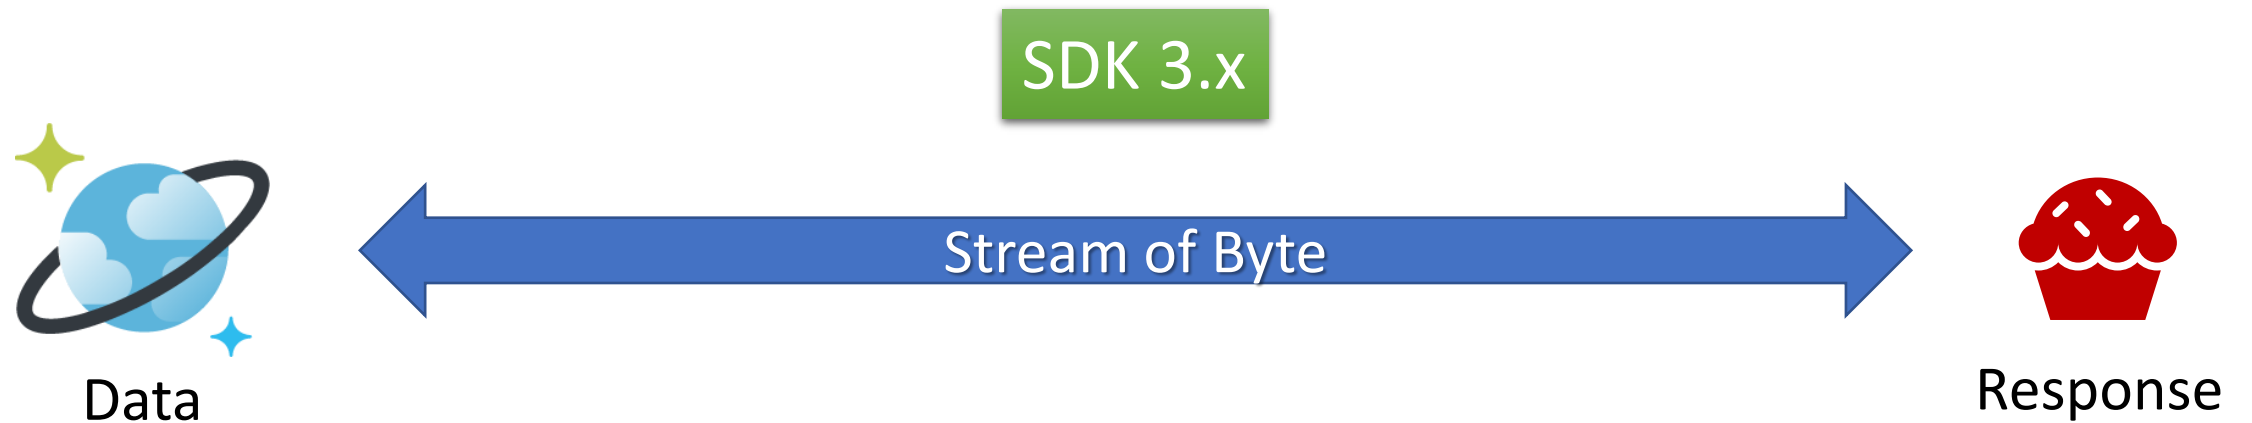
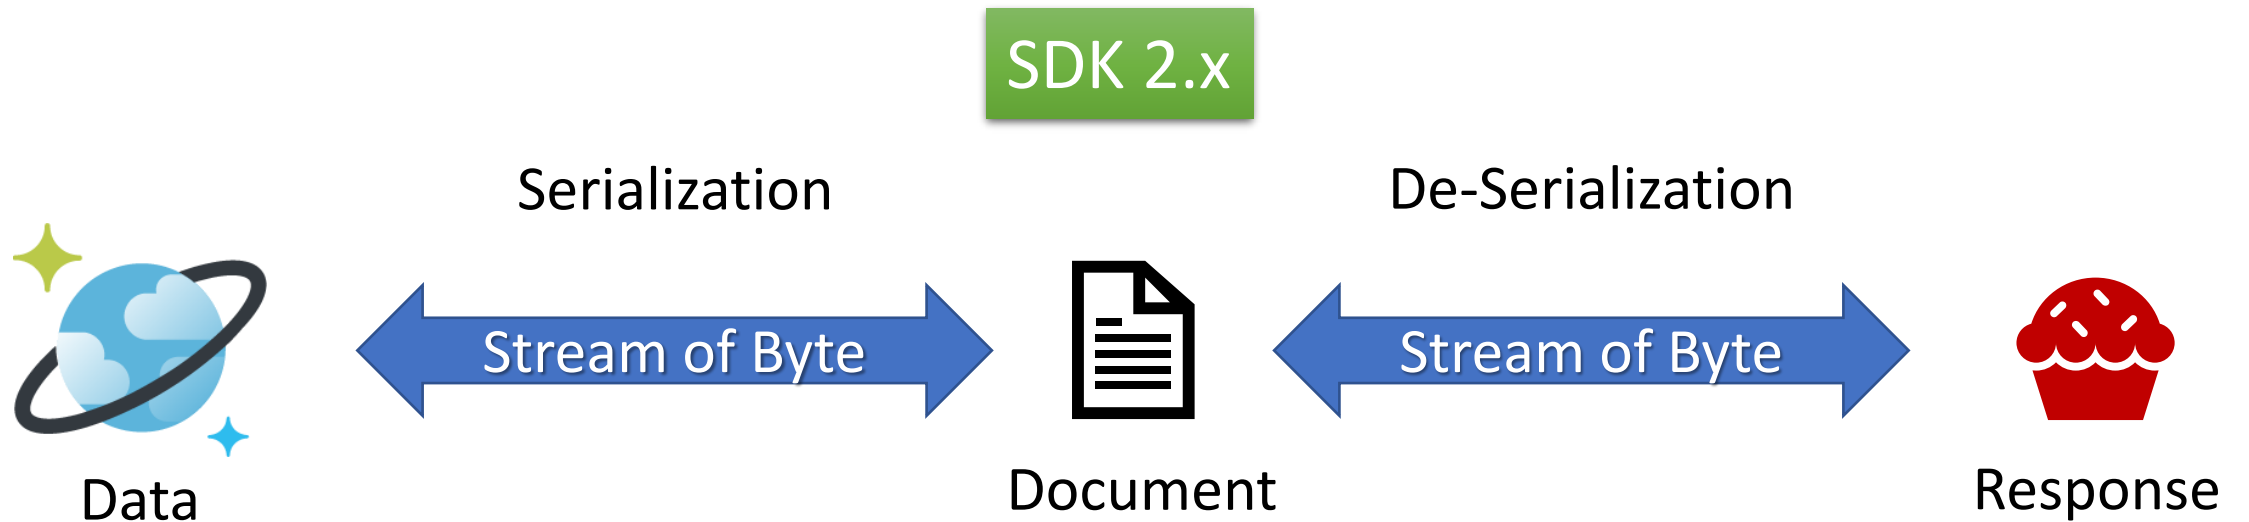
```
{  
  "runtimeOptions": {  
    "configProperties": {  
      "System.GC.Server": true  
    }  
  }  
}
```

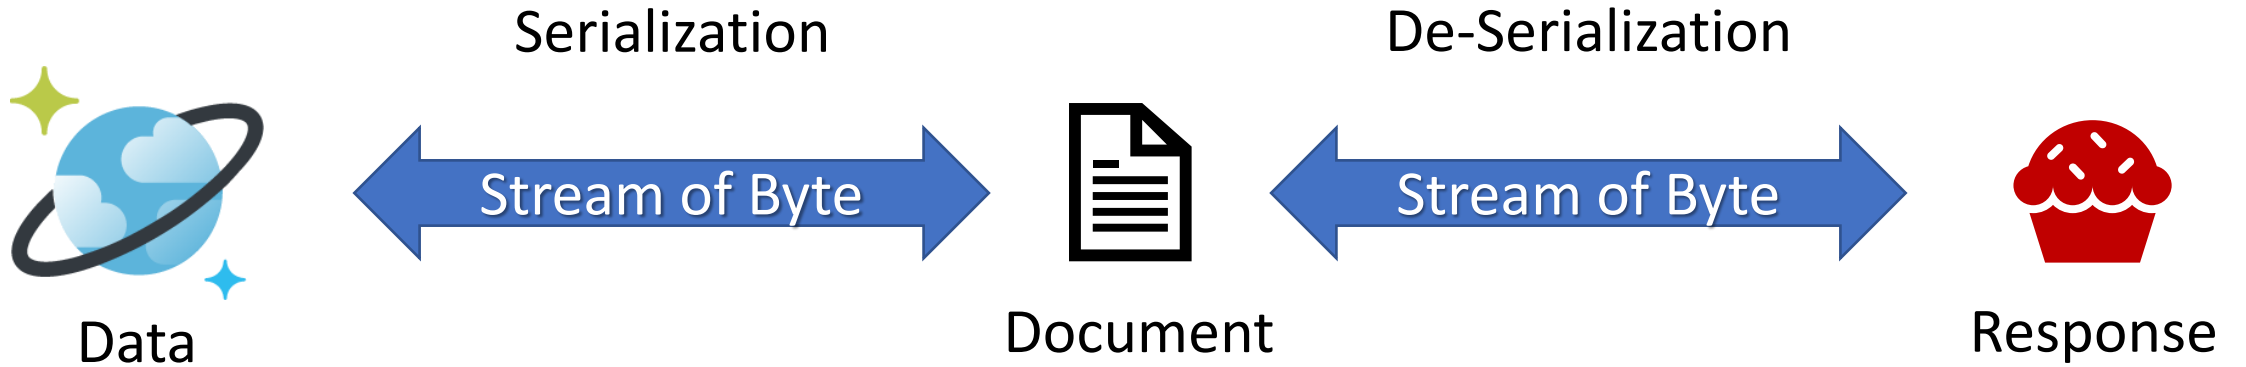
High CPU can cause increased latency and request timeout exceptions.



50.000 R/U

# Use STREAM API in SDK 3





```
public async Task RunQuery(Container container, string sqlQueryText, string partitionKeyValue)
{
    var queryDefinition = new QueryDefinition(sqlQueryText);
    FeedIterator<Sale> queryResultSetIterator = container.GetItemQueryIterator<Sale>(queryDefinition);
    var sales = new List<Sale>();
    while (queryResultSetIterator.HasMoreResults)
    {
        var currentResultSet = await queryResultSetIterator.ReadNextAsync();
        foreach (Sale current in currentResultSet)
        {
            sales.Add(current);
        }
    }
}
```



SDK 3.x



Data



Response

```
public async Task RunStreamQuery(Container container, string sqlQueryText, string partitionKeyValue)
{
    var queryDefinition = new QueryDefinition(sqlQueryText);
    FeedIterator queryResultSetIterator = container.GetItemQueryStreamIterator(queryDefinition, null,
        new QueryRequestOptions() { PartitionKey = new PartitionKey(partitionKeyValue) });
    while (queryResultSetIterator.HasMoreResults)
    {
        using (ResponseMessage response = await queryResultSetIterator.ReadNextAsync())
        {
            using (StreamReader sr = new StreamReader(response.Content))
            using (JsonTextReader jtr = new JsonTextReader(sr))
            {
                JObject result = JObject.Load(jtr);
            }
        }
    }
}
```

# DocumentClient

```
public void ConfigureServices(IServiceCollection services)
{
    services.AddControllersWithViews();

    services.AddSingleton<IDocumentClient>(
        x => new DocumentClient(new Uri("https://localhost:8081"),
            new NetworkCredential("", "masterKey").SecurePassword));
}
```

```
private readonly IDocumentClient _documentclient;
private readonly ILogger<HomeController> _logger;

0 references | 0 exceptions
public HomeController(ILogger<HomeController> logger, IDocumentClient docclient)
{
    _logger = logger;
    _documentclient = docclient;
}
```




1



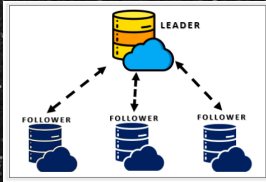
Account  
Information

2

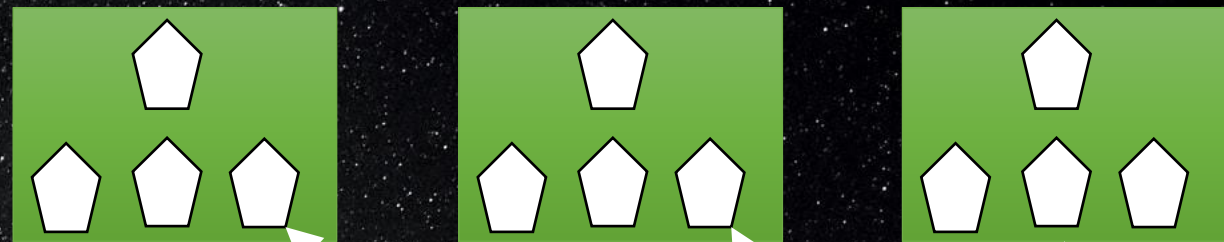


Container  
Information


3



Partition  
Information



Consistency  
Mode?



DOCUMENTCLIENT

# Optimize bandwidth

```
static async Task<bool> PushToCosmos(List<Hurricane> hurricane)
{
    try
    {
        foreach (var h in hurricane)
        {
            var temp = await client.GetDatabase("Spatial").GetContainer("Hurricanes")
                .UpsertItemAsync(h,
                    requestOptions :new ItemRequestOptions() { EnableContentResponseOnWrite=false} );
        }
        return true;
    }
    catch(Exception ex)
    {
        return false;
    }
}
```



# Partial Document Update

- Combine multiple operations
- Conditional Updates with SQL-like filters
- Supports Transactions with Transactional batch
- Conflict Resolution Support

## PARTIAL DOCUMENT OPERATIONS

Add	Creates a new element
Set	Updates an element or creates one if it does not exist.
Replace	Updates an element only if exists.
Remove	Deletes an existing element
Increment	Increases or decreases by specified value

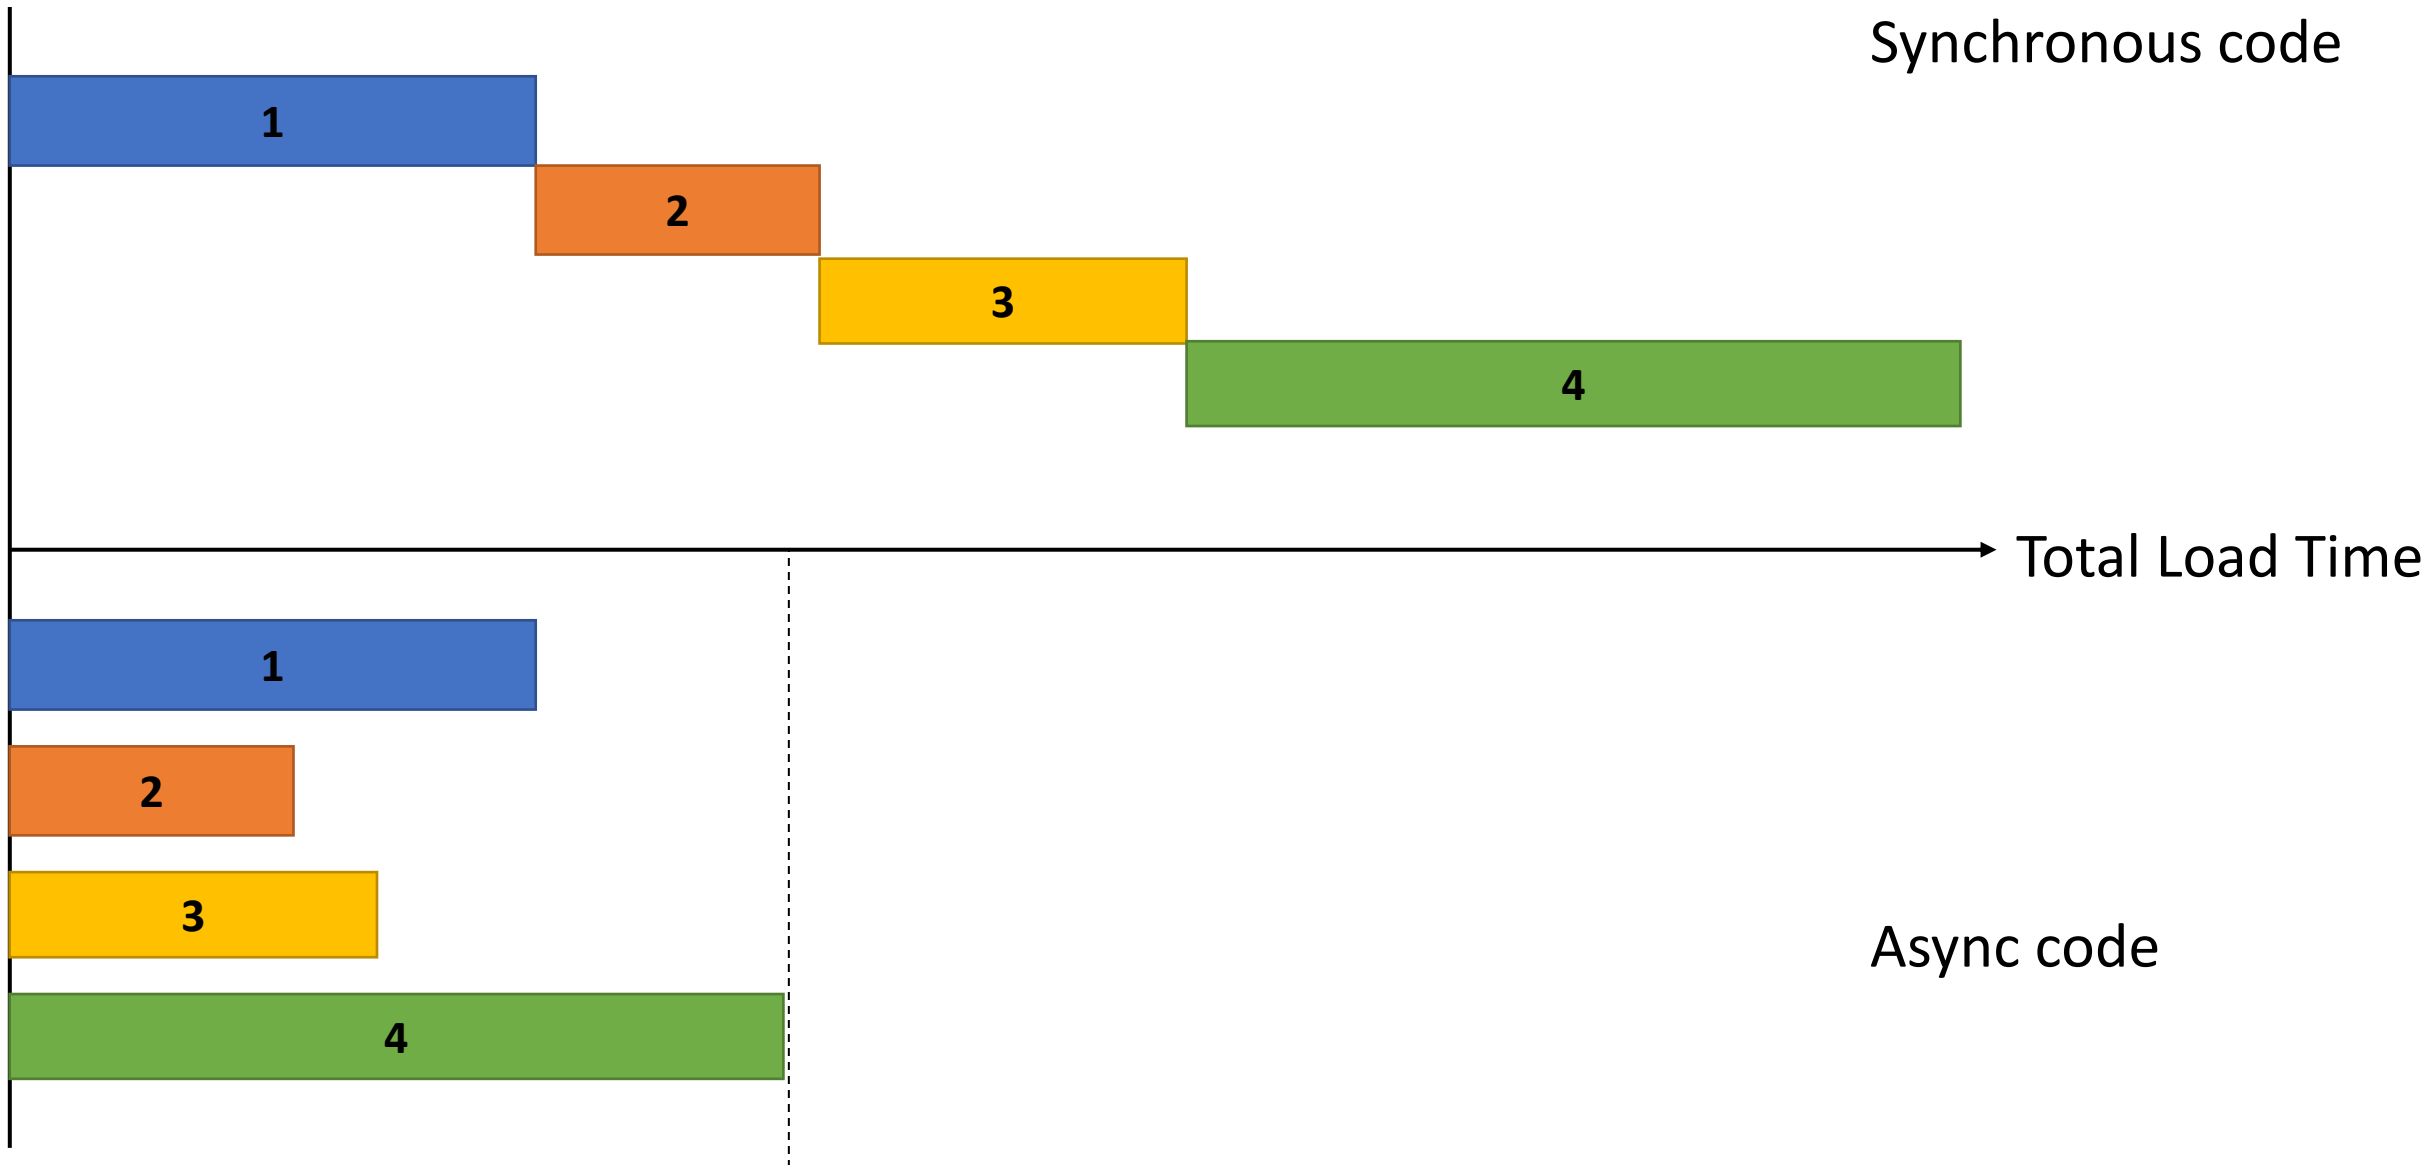
0 references | 0 requests | 0 exceptions

```
public async void PartialUpdate(Container container, string partitionKey)
{
    var ops = new List<PatchOperation>
    {
        PatchOperation.Add("/SaleDt", DateTime.UtcNow),
        PatchOperation.Remove("/OriginalSaleDt")
    };
    await container.PatchItemAsync<Sale>(
        id: "20",
        partitionKey: new PartitionKey(partitionKey),
        patchOperations: ops);
}
```

0 references | 0 requests | 0 exceptions

```
public async void PartialUpdateWithFilter(Container container, string partitionKey)
{
    var options = new PatchItemRequestOptions { FilterPredicate = "from c where c.NumberOfItems > 5" };
    await container.PatchItemAsync<Sale>(
        id: "20",
        partitionKey: new PartitionKey(partitionKey),
        patchOperations: new[] { PatchOperation.Add("/Audit", true) },
        options);
}
```

# Avoid Blocking Calls



# Indexing Policies

Available Indexing Modes

**CONSISTENT**



**LAZY**

**NONE**

COSMOS DB INDEXES EVERYTHING

Data File



Index File



Consistency  
Level





# Index Types

## RANGE INDEX

- Equality Queries
- Equality matches
- Range Queries
- String system functions
- ORDER BY
- JOIN

## SPATIAL INDEX

- Spatial functions
- Point
- LineString
- Polygon
- MultiPolygon

## COMPOSITE INDEX

- ORDER BY multiple properties
- Filter and ORDER BY
- Filter on multiple properties

# Composite Indexes

- Order BY clause with 2 or more properties requires Composite Index
- Optimize queries that have filters on multiple properties

A Composite Index can only optimize a **single** range filter or system function.

All Composite Indexed properties must be in the query's filter.

Property with a range filter must be defined **last** in the composite index.

Range Filters

Age > 10	Age < 10	Age >= 10	Age <= 10	Age != 10
----------	----------	-----------	-----------	-----------

# Modify Indexes

Exclude unused paths will

- Improve write performance
- Reduce R/U charges on write operations
- Reduce index storage

Do all indexing changes in one policy modification.

Do not use Lazy Indexing Mode

Take advantage of Composite Indexes

Indexing Policy

```
1  {
2      "indexingMode": "consistent",
3      "automatic": true,
4      "includedPaths": [
5          {
6              "path": "/*"
7          }
8      ],
9      "excludedPaths": [
10         {
11             "path": "/\"_etag\"/?"
12         }
13     ],
```



## Provisioned



## Auto Scale



## Serverless

Daily used Applications

Predictable/Stable Workloads

Any Traffic

All Types of Applications

Production friendly

Minimum: 400 RU/s

Infrequently Used Applications

Unpredictable/Critical Workloads

Any Traffic

New Applications

Dev & Test Environments

Minimum: 400 – 4000 RU/s

Infrequently Used Applications

Limited Workloads

Light Traffic / Long Idle times

Proof of Concept Applications

Dev & Test Environments

No Global Distribution

5000 RU/s and 50 GB per container

**\$0.008/hour for 100 RU/s**

**\$0.008/hour for 100 RU/s**

**\$0.25 per 1M RU**

# Execution Metrics

```
var query = client.CreateDocumentQuery(
    UriFactory.CreateDocumentCollectionUri(_dbname, _container),
    $"SELECT * FROM Orders o WHERE o.OrderId = {@orderid}",
    new FeedOptions {
        PopulateQueryMetrics = true,
        EnableCrossPartitionQuery = true }).AsDocumentQuery();
FeedResponse<dynamic> result = await query.ExecuteNextAsync();
IReadOnlyDictionary<string, QueryMetrics> metrics = result.QueryMetrics;
```

metrics	Count = 1
[0]	{[0, Retrieved Document Count : 0 ...
Key	"0"
Value	{Retrieved Document Count : 0 ...
IndexHitRatio	1
OutputDocumentCount	0
QueryEngineTimes	{Microsoft.Azure.Documents.QueryEngineTimes}
DocumentLoadTime	{00:00:00}
IndexLookupTime	{00:00:00.0000700}
RuntimeExecutionTimes	{Microsoft.Azure.Documents.RuntimeExecutionTimes}
SystemFunctionExecutionTime	{00:00:00}
TotalTime	{00:00:00.0000200}
UserDefinedFunctionExecutionTime	{00:00:00}
Static members	
Non-Public members	
WriteOutputTime	{00:00:00}
Non-Public members	
QueryPreparationTimes	{Microsoft.Azure.Documents.QueryPreparationTimes}
CompileTime	{00:00:00.0000700}
LogicalPlanBuildTime	{00:00:00.0000300}
PhysicalPlanBuildTime	{00:00:00.0000600}
QueryOptimizationTime	{00:00:00}
Static members	
Non-Public members	
Retries	0
RetrievedDocumentCount	0
RetrievedDocumentSize	0
TotalTime	{00:00:00.0004900}

# Thank you for joining us

Future Data Driven 2022

## Scan the QR Code

Join The Data Driven Community





# Thank you!

You can reach me on the following platforms.



SavranWeb



hasansavran



<https://h-savran.blogspot.com/>